

[Start Here](#)[Blog](#)[Books](#)[About](#)[Contact](#)

# How to Make Out-of-Sample Forecasts with ARIMA in Python

by **Jason Brownlee** on March 24, 2017 in **Time Series**



Making out-of-sample forecasts can be confusing when getting started with time series data.

The statsmodels Python API provides functions for performing one-step and multi-step out-of-sample forecasts.

In this tutorial, you will clear up any confusion you have about making out-of-sample forecasts with time series data in Python.

After completing this tutorial, you will know:

- How to make a one-step out-of-sample forecast.
- How to make a multi-step out-of-sample forecast.
- The difference between the *forecast()* and *predict()* functions.

Let's get started.



How to Make Out-of-Sample Forecasts with ARIMA in Python  
Photo by [dziambel](#), some rights reserved.

## Tutorial Overview

This tutorial is broken down into the following 5 steps:

1. Dataset Description
2. Split Dataset
3. Develop Model
4. One-Step Out-of-Sample Forecast
5. Multi-Step Out-of-Sample Forecast

---

### Stop learning Time Series Forecasting the *slow* way!

Take my free 7-day email course and discover data prep, modeling and more (with sample code).

Click to sign-up and also get a free PDF Ebook version of the course.

**Start Your FREE Mini-Course Now!**

---

## 1. Minimum Daily Temperatures Dataset

This dataset describes the minimum daily temperatures over 10 years (1981-1990) in the city of Melbourne, Australia.

The units are in degrees Celsius and there are 3,650 observations. The source of the data is credited as the Australian Bureau of Meteorology.

[Learn more about the dataset on Data Market.](#)

Download the Minimum Daily Temperatures dataset to your current working directory with the filename *"daily-minimum-temperatures.csv"*.

**Note:** The downloaded file contains some question mark ("?) characters that must be removed before you can use the dataset. Open the file in a text editor and remove the "?" characters. Also, remove any footer information in the file.


The example below loads the dataset as a Pandas Series.

```
1 # line plot of time series
2 from pandas import Series
3 from matplotlib import pyplot
4 # load dataset
5 series = Series.from_csv('daily-minimum-temperatures.csv', header=0)
6 # display first few rows
7 print(series.head(20))
8 # line plot of dataset
9 series.plot()
10 pyplot.show()
```

Running the example prints the first 20 rows of the loaded dataset.

```
1 Date
2 1981-01-01    20.7
3 1981-01-02    17.9
4 1981-01-03    18.8
5 1981-01-04    14.6
6 1981-01-05    15.8
7 1981-01-06    15.8
8 1981-01-07    15.8
9 1981-01-08    17.4
10 1981-01-09    21.8
11 1981-01-10    20.0
12 1981-01-11    16.2
13 1981-01-12    13.3
14 1981-01-13    16.7
15 1981-01-14    21.5
16 1981-01-15    25.0
17 1981-01-16    20.7
18 1981-01-17    20.6
19 1981-01-18    24.8
20 1981-01-19    17.7
21 1981-01-20    15.5
```

A line plot of the time series is also created.



Minimum Daily Temperatures Dataset Line Plot

## 2. Split Dataset

We can split the dataset into two parts.

The first part is the training dataset that we will use to prepare an ARIMA model. The second part is the test dataset that we will pretend is not available. It is these time steps that we will treat as out of sample.

The dataset contains data from January 1st 1981 to December 31st 1990.

We will hold back the last 7 days of the dataset from December 1990 as the test dataset and treat those time steps as out of sample.

Specifically 1990-12-25 to 1990-12-31:

```
1 1990-12-25,12.9
2 1990-12-26,14.6
3 1990-12-27,14.0
4 1990-12-28,13.6
5 1990-12-29,13.5
6 1990-12-30,15.7
7 1990-12-31,13.0
```

The code below will load the dataset, split it into the training and validation datasets, and save them to files *dataset.csv* and *validation.csv* respectively.

```
1 # split the dataset
2 from pandas import Series
```

```

3 series = Series.from_csv('daily-minimum-temperatures.csv', header=0)
4 split_point = len(series) - 7
5 dataset, validation = series[0:split_point], series[split_point:]
6 print('Dataset %d, Validation %d' % (len(dataset), len(validation)))
7 dataset.to_csv('dataset.csv')
8 validation.to_csv('validation.csv')

```

Run the example and you should now have two files to work with.

The last observation in the dataset.csv is Christmas Eve 1990:

```
1 1990-12-24,10.0
```

That means Christmas Day 1990 and onwards are out-of-sample time steps for a model trained on *dataset.csv*.

### 3. Develop Model

In this section, we are going to make the data stationary and develop a simple ARIMA model.

The data has a strong seasonal component. We can neutralize this and make the data stationary by taking the seasonal difference. That is, we can take the observation for a day and subtract the observation from the same day one year ago.

This will result in a stationary dataset from which we can fit a model.

```

1 # create a differenced series
2 def difference(dataset, interval=1):
3     diff = list()
4     for i in range(interval, len(dataset)):
5         value = dataset[i] - dataset[i - interval]
6         diff.append(value)
7     return numpy.array(diff)

```

We can invert this operation by adding the value of the observation one year ago. We will need to do this to any forecasts made by a model trained on the seasonally adjusted data.

```

1 # invert differenced value
2 def inverse_difference(history, yhat, interval=1):
3     return yhat + history[-interval]

```

We can fit an ARIMA model.

Fitting a strong ARIMA model to the data is not the focus of this post, so rather than going through the analysis of the problem or [grid searching parameters](#), I will choose a simple ARIMA(7,0,7) configuration.

We can put all of this together as follows:

```

1 from pandas import Series
2 from statsmodels.tsa.arima_model import ARIMA
3 import numpy
4
5 # create a differenced series
6 def difference(dataset, interval=1):
7     diff = list()
8     for i in range(interval, len(dataset)):
9         value = dataset[i] - dataset[i - interval]
10        diff.append(value)

```

```

11     return numpy.array(diff)
12
13 # load dataset
14 series = Series.from_csv('dataset.csv', header=None)
15 # seasonal difference
16 X = series.values
17 days_in_year = 365
18 differenced = difference(X, days_in_year)
19 # fit model
20 model = ARIMA(differenced, order=(7,0,1))
21 model_fit = model.fit(dis=0)
22 # print summary of fit model
23 print(model_fit.summary())

```

Running the example loads the dataset, takes the seasonal difference, then fits an ARIMA(7,0,7) model and prints the summary of the fit model.

```

1                      ARMA Model Results
2  =====
3 Dep. Variable:          y      No. Observations:          3278
4 Model:                ARMA(7, 1)  Log Likelihood          -8673.748
5 Method:                css-mle    S.D. of innovations      3.411
6 Date:                  Mon, 20 Feb 2017  AIC                17367.497
7 Time:                  10:28:38    BIC                    17428.447
8 Sample:                0          HQIC                   17389.322
9
10  =====
11      coef      std err          z      P>|z|      [0.025      0.975]
12  -----
13 const          0.0132      0.132      0.100      0.921      -0.246      0.273
14 ar.L1.y         1.1424      0.287      3.976      0.000      0.579      1.706
15 ar.L2.y        -0.4346      0.154     -2.829      0.005     -0.736     -0.133
16 ar.L3.y         0.0961      0.042      2.289      0.022      0.014      0.178
17 ar.L4.y         0.0125      0.029      0.434      0.664     -0.044      0.069
18 ar.L5.y        -0.0101      0.029     -0.343      0.732     -0.068      0.047
19 ar.L6.y         0.0119      0.027      0.448      0.654     -0.040      0.064
20 ar.L7.y         0.0089      0.024      0.368      0.713     -0.038      0.056
21 ma.L1.y        -0.6157      0.287     -2.146      0.032     -1.178     -0.053
22
23                      Roots
24  =====
25      Real      Imaginary      Modulus      Frequency
26  -----
27 AR.1          1.2234          -0.0000j          1.2234          -0.0000
28 AR.2          1.2561          -1.0676j          1.6485          -0.1121
29 AR.3          1.2561          +1.0676j          1.6485           0.1121
30 AR.4          0.0349          -2.0160j          2.0163          -0.2472
31 AR.5          0.0349          +2.0160j          2.0163           0.2472
32 AR.6         -2.5770          -1.3110j          2.8913          -0.4251
33 AR.7         -2.5770          +1.3110j          2.8913           0.4251
34 MA.1          1.6242          +0.0000j          1.6242           0.0000

```

We are now ready to explore making out-of-sample forecasts with the model.

## 4. One-Step Out-of-Sample Forecast

ARIMA models are great for one-step forecasts.

A one-step forecast is a forecast of the very next time step in the sequence from the available data used to fit the model.

In this case, we are interested in a one-step forecast of Christmas Day 1990:

```
1 1990-12-25
```

## Forecast Function

The statsmodel ARIMAResults object provides a *forecast()* function for making predictions.

By default, this function makes a single step out-of-sample forecast. As such, we can call it directly and make our forecast. The result of the *forecast()* function is an array containing the forecast value, the standard error of the forecast, and the confidence interval information. Now, we are only interested in the first element of this forecast, as follows.

```
1 # one-step out-of sample forecast
2 forecast = model_fit.forecast()[0]
```

Once made, we can invert the seasonal difference and convert the value back into the original scale.

```
1 # invert the differenced forecast to something usable
2 forecast = inverse_difference(X, forecast, days_in_year)
```

The complete example is listed below:

```
1 from pandas import Series
2 from statsmodels.tsa.arima_model import ARIMA
3 import numpy
4
5 # create a differenced series
6 def difference(dataset, interval=1):
7     diff = list()
8     for i in range(interval, len(dataset)):
9         value = dataset[i] - dataset[i - interval]
10        diff.append(value)
11    return numpy.array(diff)
12
13 # invert differenced value
14 def inverse_difference(history, yhat, interval=1):
15     return yhat + history[-interval]
16
17 # load dataset
18 series = Series.from_csv('dataset.csv', header=None)
19 # seasonal difference
20 X = series.values
21 days_in_year = 365
22 differenced = difference(X, days_in_year)
23 # fit model
24 model = ARIMA(differenced, order=(7,0,1))
25 model_fit = model.fit(disp=0)
26 # one-step out-of sample forecast
27 forecast = model_fit.forecast()[0]
28 # invert the differenced forecast to something usable
29 forecast = inverse_difference(X, forecast, days_in_year)
30 print('Forecast: %f' % forecast)
```

Running the example prints 14.8 degrees, which is close to the expected 12.9 degrees in the *validation.csv* file.

```
1 Forecast: 14.861669
```

## Predict Function

The statsmodel ARIMAResults object also provides a *predict()* function for making forecasts.

The predict function can be used to predict arbitrary in-sample and out-of-sample time steps, including the next out-of-sample forecast time step.



The predict function requires a start and an end to be specified, these can be the indexes of the time steps relative to the beginning of the training data used to fit the model, for example:

```
1 # one-step out of sample forecast
2 start_index = len(differenced)
3 end_index = len(differenced)
4 forecast = model_fit.predict(start=start_index, end=end_index)
```

The start and end can also be a datetime string or a “datetime” type; for example:

```
1 start_index = '1990-12-25'
2 end_index = '1990-12-25'
3 forecast = model_fit.predict(start=start_index, end=end_index)
```

and

```
1 from pandas import datetime
2 start_index = datetime(1990, 12, 25)
3 end_index = datetime(1990, 12, 26)
4 forecast = model_fit.predict(start=start_index, end=end_index)
```

Using anything other than the time step indexes results in an error on my system, as follows:

```
1 AttributeError: 'NoneType' object has no attribute 'get_loc'
```

Perhaps you will have more luck; for now, I am sticking with the time step indexes.

The complete example is listed below:

```
1 from pandas import Series
2 from statsmodels.tsa.arima_model import ARIMA
3 import numpy
4 from pandas import datetime
5
6 # create a differenced series
7 def difference(dataset, interval=1):
8     diff = list()
9     for i in range(interval, len(dataset)):
10         value = dataset[i] - dataset[i - interval]
11         diff.append(value)
12     return numpy.array(diff)
13
14 # invert differenced value
15 def inverse_difference(history, yhat, interval=1):
16     return yhat + history[-interval]
17
18 # load dataset
19 series = Series.from_csv('dataset.csv', header=None)
20 # seasonal difference
21 X = series.values
22 days_in_year = 365
23 differenced = difference(X, days_in_year)
24 # fit model
25 model = ARIMA(differenced, order=(7,0,1))
26 model_fit = model.fit(disp=0)
27 # one-step out of sample forecast
28 start_index = len(differenced)
29 end_index = len(differenced)
30 forecast = model_fit.predict(start=start_index, end=end_index)
31 # invert the differenced forecast to something usable
32 forecast = inverse_difference(X, forecast, days_in_year)
33 print('Forecast: %f' % forecast)
```

Running the example prints the same forecast as above when using the *forecast()* function.



```
1 Forecast: 14.861669
```

You can see that the predict function is more flexible. You can specify any point or contiguous forecast interval in or out of sample.

Now that we know how to make a one-step forecast, we can now make some multi-step forecasts.

## 5. Multi-Step Out-of-Sample Forecast

We can also make multi-step forecasts using the *forecast()* and *predict()* functions.

It is common with weather data to make one week (7-day) forecasts, so in this section we will look at predicting the minimum daily temperature for the next 7 out-of-sample time steps.

### Forecast Function

The *forecast()* function has an argument called *steps* that allows you to specify the number of time steps to forecast.

By default, this argument is set to 1 for a one-step out-of-sample forecast. We can set it to 7 to get a forecast for the next 7 days.

```
1 # multi-step out-of-sample forecast
2 forecast = model_fit.forecast(steps=7)[0]
```

We can then invert each forecasted time step, one at a time and print the values. Note that to invert the forecast value for  $t+2$ , we need the inverted forecast value for  $t+1$ . Here, we add them to the end of a list called *history* for use when calling *inverse\_difference()*.

```
1 # invert the differenced forecast to something usable
2 history = [x for x in X]
3 day = 1
4 for yhat in forecast:
5     inverted = inverse_difference(history, yhat, days_in_year)
6     print('Day %d: %f' % (day, inverted))
7     history.append(inverted)
8     day += 1
```

The complete example is listed below:

```
1 from pandas import Series
2 from statsmodels.tsa.arima_model import ARIMA
3 import numpy
4
5 # create a differenced series
6 def difference(dataset, interval=1):
7     diff = list()
8     for i in range(interval, len(dataset)):
9         value = dataset[i] - dataset[i - interval]
10        diff.append(value)
11    return numpy.array(diff)
12
13 # invert differenced value
14 def inverse_difference(history, yhat, interval=1):
15     return yhat + history[-interval]
16
17 # load dataset
18 series = Series.from_csv('dataset.csv', header=None)
19 # seasonal difference
```

```

20 X = series.values
21 days_in_year = 365
22 differenced = difference(X, days_in_year)
23 # fit model
24 model = ARIMA(differenced, order=(7,0,1))
25 model_fit = model.fit(dispatch=0)
26 # multi-step out-of-sample forecast
27 forecast = model_fit.forecast(steps=7)[0]
28 # invert the differenced forecast to something usable
29 history = [x for x in X]
30 day = 1
31 for yhat in forecast:
32     inverted = inverse_difference(history, yhat, days_in_year)
33     print('Day %d: %f' % (day, inverted))
34     history.append(inverted)
35     day += 1

```

Running the example prints the forecast for the next 7 days.

```

1 Day 1: 14.861669
2 Day 2: 15.628784
3 Day 3: 13.331349
4 Day 4: 11.722413
5 Day 5: 10.421523
6 Day 6: 14.415549
7 Day 7: 12.674711

```

## Predict Function

The *predict()* function can also forecast the next 7 out-of-sample time steps.

Using time step indexes, we can specify the end index as 6 more time steps in the future; for example:

```

1 # multi-step out-of-sample forecast
2 start_index = len(differenced)
3 end_index = start_index + 6
4 forecast = model_fit.predict(start=start_index, end=end_index)

```

The complete example is listed below.

```

1 from pandas import Series
2 from statsmodels.tsa.arima_model import ARIMA
3 import numpy
4
5 # create a differenced series
6 def difference(dataset, interval=1):
7     diff = list()
8     for i in range(interval, len(dataset)):
9         value = dataset[i] - dataset[i - interval]
10        diff.append(value)
11    return numpy.array(diff)
12
13 # invert differenced value
14 def inverse_difference(history, yhat, interval=1):
15     return yhat + history[-interval]
16
17 # load dataset
18 series = Series.from_csv('dataset.csv', header=None)
19 # seasonal difference
20 X = series.values
21 days_in_year = 365
22 differenced = difference(X, days_in_year)
23 # fit model
24 model = ARIMA(differenced, order=(7,0,1))
25 model_fit = model.fit(dispatch=0)
26 # multi-step out-of-sample forecast

```

```
27 start_index = len(differenced)
28 end_index = start_index + 6
29 forecast = model_fit.predict(start=start_index, end=end_index)
30 # invert the differenced forecast to something usable
31 history = [x for x in X]
32 day = 1
33 for yhat in forecast:
34     inverted = inverse_difference(history, yhat, days_in_year)
35     print('Day %d: %f' % (day, inverted))
36     history.append(inverted)
37     day += 1
```

Running the example produces the same results as calling the *forecast()* function in the previous section, as you would expect.

```
1 Day 1: 14.861669
2 Day 2: 15.628784
3 Day 3: 13.331349
4 Day 4: 11.722413
5 Day 5: 10.421523
6 Day 6: 14.415549
7 Day 7: 12.674711
```

## Summary

In this tutorial, you discovered how to make out-of-sample forecasts in Python using statsmodels.


Specifically, you learned:

- How to make a one-step out-of-sample forecast.
- How to make a 7-day multi-step out-of-sample forecast.
- How to use both the *forecast()* and *predict()* functions when forecasting.

Do you have any questions about out-of-sample forecasts, or about this post? Ask your questions in the comments and I will do my best to answer.

---

## Want to Develop Time Series Forecasts with Python?



### Develop Your Own Forecasts in Minutes

...with just a few lines of python code

Discover how in my new Ebook:

[Introduction to Time Series Forecasting With Python](#)

It covers **self-study tutorials** and **end-to-end projects** on topics like:  
*Loading data, visualization, modeling, algorithm tuning, and much more...*

### Finally Bring Time Series Forecasting to Your Own Projects

Skip the Academics. Just Results.

[Click to learn more.](#)



### About Jason Brownlee

Dr. Jason Brownlee is a husband, proud father, academic researcher, author, professional developer and a machine learning practitioner. He is dedicated to helping developers get started and get good at applied machine learning.

[Learn more.](#)

[View all posts by Jason Brownlee →](#)

< Time Series Forecasting with Python 7-Day Mini-Course

Sensitivity Analysis of History Size to Forecast Skill with ARIMA in Python >

## 55 Responses to *How to Make Out-of-Sample Forecasts with ARIMA in Python*



**Steve** March 24, 2017 at 10:44 pm #

REPLY ↩

Your tutorials are the most helpful machine learning resources I have found on the Internet and have been hugely helpful in work and personal side projects. I don't know if you take requests but I'd love to see a series of posts on recommender systems one of these days!



**Jason Brownlee** March 25, 2017 at 7:36 am #

REPLY ↩

Thanks Steve, and great suggestion! Thanks.



**Tim** April 27, 2017 at 12:43 pm #

REPLY ↩

Hi,

This is a really nice example. Do you know if the ARIMA class allows to define the specification of the model without going through the fitting procedure. Let's say I have parameters that were estimated using a dataset that I no longer have but I still want to produce a forecast.

Thanks



**Jason Brownlee** April 28, 2017 at 7:32 am #

REPLY ↩

I expect you can set the coefficients explicitly within the ARIMA model.

Sorry I do not have an example, this post may be relevant:

<http://machinelearningmastery.com/make-manual-predictions-arma-models-python/>



**masum** May 11, 2017 at 8:32 pm #

REPLY ↩

sir,

would it be possible to do the same using LSTM RNN ?

if it is would you please come up with a blog?

Thanking you



**Jason Brownlee** May 12, 2017 at 7:41 am #

REPLY ↩

Yes.

Any of my LSTM tutorials show how to make out of sample forecasts. For example:

<http://machinelearningmastery.com/multi-step-time-series-forecasting-long-short-term-memory-networks-python/>



**masum** May 12, 2017 at 8:29 pm #

REPLY ↩

I tried to run the above example without any seasonal difference with given below code.

```
from pandas import Series
from matplotlib import pyplot
from pandas import Series
from statsmodels.tsa.arma_model import ARIMA
# load dataset
series = Series.from_csv('daily-minimum-temperatures.csv', header=0)
print(series.head(20))
series.plot()
pyplot.show()

split_point = len(series) - 7
dataset, validation = series[0:split_point], series[split_point:]
print('Dataset %d, Validation %d' % (len(dataset), len(validation)))
dataset.to_csv('dataset.csv')
validation.to_csv('validation.csv')

series = Series.from_csv('dataset.csv', header=None)
model = ARIMA(series, order=(7,0,1))
model_fit = model.fit(dispatch=0)
```

```
forecast = model_fit.forecast(steps=7)[0]
print('Forecast: %f' % forecast)
```

for the code i am getting an error:

TypeError: only length-1 arrays can be converted to Python scalars

how can i solve this? it does well for single step forecast



**Jason Brownlee** May 13, 2017 at 6:13 am #

REPLY ↩

I would recommend double checking your data, make sure any footer information was deleted.



**Hans** June 1, 2017 at 12:58 am #

REPLY ↩

What does 'seasonal difference' mean?

And what are the details of:

'Once made, we can invert the seasonal difference and convert the value back into the original scale.'

Is it worth to test this code with non-seasonal data or is there another ARIMA-tutorial for non-seasonal approaches on this site?



**Jason Brownlee** June 2, 2017 at 12:51 pm #

REPLY ↩

See this post:

<http://machinelearningmastery.com/seasonal-persistence-forecasting-python/>

And this post:

<http://machinelearningmastery.com/time-series-seasonality-with-python/>

Please use the search feature of the blog.



**Hans** June 15, 2017 at 11:27 am #

REPLY ↩

If I pretend data in test-partition is not given, does this tutorial do the same except of the seasonal cleaning?

<http://machinelearningmastery.com/tune-arima-parameters-python/>



**Hans** June 15, 2017 at 11:29 am #

REPLY ↩

Can I obtain a train RMSE from this example. Is training involved?



**Jason Brownlee** June 16, 2017 at 7:47 am #

REPLY ↩

The model is trained, then the trained model is used to make a forecast.

Consider reading and working through the tutorial.



**Hans** June 16, 2017 at 12:16 pm #

REPLY ↩

I did so several times.

How can I obtain a train RMSE from the model?



**Jason Brownlee** June 17, 2017 at 7:20 am #

REPLY ↩

See this post on how to estimate the skill of a model prior to using it to make out of sample predictions:

<http://machinelearningmastery.com/backtest-machine-learning-models-time-series-forecasting/>

See this post to understand the difference between evaluating a model and using a final model to make predictions:

<http://machinelearningmastery.com/train-final-machine-learning-model/>



**Hans** June 19, 2017 at 5:35 am #

I actually meant obtain a train RMSE from the model in the example.  
As I understand the model was trained before making an out of sample prediction.  
If we place a

```
print(model_fit.summary())
```

right after fitting/training it prints some information's, but no train RMSE.

A)

Is there a way to use the summary-information to obtain a train RMSE?

B)

Is there a way in Python to obtain all properties and methods from the model\_fit object-like in other languages?



**Jason Brownlee** June 19, 2017 at 8:47 am #

Yes, this tutorial assumes you have already estimated the skill of your model and are now ready to use it to make forecasts.



Estimating the skill of the model is a different task. You can do this using walk forward validation or a train/test split evaluation.



**Hans** June 16, 2017 at 3:06 pm #

REPLY ↩

Is this the line where the training happens?

```
model = ARIMA(differenced, order=(7,0,1))
```



**Jason Brownlee** June 17, 2017 at 7:22 am #

REPLY ↩

No here:

```
1 model_fit = model.fit(dispatch=0)
```



**Hans** June 25, 2017 at 12:29 pm #

REPLY ↩

Yes I know. I actually thought there could be a direct answer to A) and B).  
I would use it for archiving.



**Hans** June 15, 2017 at 12:40 pm #

REPLY ↩

If I write: 'split\_point = len(series) - 0' while my last datapoint in dataset is from today.

Would I have a valid forecast for tomorrow?



**M.Swefy** June 22, 2017 at 12:39 am #

REPLY ↩

thanks a lot for the nice detailed article, i followed all steps and they all seem working properly, i seek your support Dr. to help me organize my project.

i have a raw data for temperature readings for some nodes (hourly readings), i selected the training set and divided them to test and training sets.

i used ARIMA model to train and test and i got Test MSE: 3.716.

now i need to expose the mass raw data to the trained model, then get the forecasted values vs. the actual values in the same csv file.

what should i do



**M.Swefy** June 22, 2017 at 12:41 am #

REPLY ↩

\*ARIMA

**Jason Brownlee** June 22, 2017 at 6:09 am #

REPLY ↩

I'm not sure I follow. Consider this post on how to evaluate a time series model:  
<http://machinelearningmastery.com/backtest-machine-learning-models-time-series-forecasting/>

**AMU** June 23, 2017 at 5:33 am #

REPLY ↩

Thank you Jason for this wonderful post... It is very detailed and easy to understand..

Do you also have something similar for LSTM Neural Network algorithm as well? something like – How to Make Out-of-Sample Forecasts with LSTM in Python.

If not, will you write one blog like this with detail explanation? I am sure there are lot of people have the same question.

**Jason Brownlee** June 23, 2017 at 6:45 am #

REPLY ↩

Almost every post I have on LSTMs shows how to make out of sample forecasts. The code is wrapped up in the walk-forward validation.

**Franklin** July 1, 2017 at 1:09 am #

REPLY ↩

Hi Jason,

Thanks a lot for this lesson. It was pretty straightforward and easy to follow. It would have been a nice bonus to show how to evaluate the forecasts though with standard metrics. We separated the validation set out and forecasted values for that week, but didn't compare to see how accurate the forecast was.

On that note, I want to ask, does it make sense to use  $R^2$  to score a time series forecast against test data? I'm trying to create absolute benchmarks for a time series that I'm analyzing and want to report unit-independent metrics, i.e. not standard RMSE that is necessarily expressed in the problem's unit scale. What about standardizing the data using zero mean and unit variance, fitting ARIMA, forecasting, and reporting that RMSE? I've been doing this and taking the  $R^2$  and the results are pretty interpretable. RMSE: 0.149 /  $R^2$ : 0.8732, but I'm just wondering if doing things this way doesn't invalidate something along the way. Just want to be correct in my process.

Thanks!

**Jason Brownlee** July 1, 2017 at 6:37 am #

REPLY ↩

We do that in other posts. Tens of other posts in fact.

This post was laser focused on “how do I make a prediction when I don’t know the real answer”.

Yes, if  $R^2$  is meaningful to you, that you can interpret it in your domain.

Generally, I recommend inverting all transforms on the prediction and then evaluating model skill at least for RMSE or MAE where you want apples-to-apples. This may be less of a concern for an  $R^2$ .



**Vishanth** July 19, 2017 at 6:56 am #

REPLY ↩

Seriously amazing. Thanks a lot professor



**Jason Brownlee** July 19, 2017 at 8:30 am #

REPLY ↩

Thanks. Also, I'm not a professor.



**Kirui** July 20, 2017 at 5:15 pm #

REPLY ↩

I get this error from your code

Traceback (most recent call last):

File “..”, line 22, in

differenced = difference(X, days\_in\_year)

File “..”, line 9, in difference

value = dataset[i] – dataset[i – interval]

TypeError: unsupported operand type(s) for -: ‘str’ and ‘str’

Cant tell where the problem is.



**Jason Brownlee** July 21, 2017 at 9:31 am #

REPLY ↩

Perhaps check that you have loaded your data correct (as real values) and that you have copied all of the code from the post without extra white space.



**Antoine** August 23, 2017 at 1:00 am #

REPLY ↩

Hi Jason,

Thanks for this detailed explanation. Very clear.

Do you know if it is possible to use the fitted parameters of an ARMA model (ARMAResults.params) and apply it on an other data set ?

I have an online process that compute a forecasting and I would like to have only one learning process (one usage of the fit() function). The rest of the time, I would like to applied the previously found parameters to the data.

Thanks in advance !



**Jason Brownlee** August 23, 2017 at 6:56 am #

REPLY ↩

Yes, you can use a grid search:

<https://machinelearningmastery.com/grid-search-arima-hyperparameters-with-python/>



**Bob** October 6, 2017 at 11:53 pm #

REPLY ↩

Ciao Jason,

Thanks for this tutorial and all the time series related ones. There is always a sense of order in how you write both posts and code.

I'm by the way still confused about something which is probably more conceptual about ARIMA.

The ARIMA parameters specify the lag which it uses to forecast.

In your case you used  $p=7$  for example so that you would take into consideration the previous week.

A first silly question is why do I need to fit an entire year of data if I'm only looking at my window/lags ?

The second question is that fitting my model I get an error which is really minimal even if I use a short training (2 days vs 1 year) which would reinforce my first point.

What am I missing?

Thanks



**Jason Brownlee** October 7, 2017 at 5:56 am #

REPLY ↩

The model needs lots of examples in order to generalize to new cases.

More data is often better, to a point of diminishing returns in terms of model skill.



**Kai** October 31, 2017 at 12:02 pm #

REPLY ↩

Hi Jason. Thanks for this awesome post.

But I have a question that is it possible to fit a multivariable time series using ARIMA model? Let's say we have a 312-dimension at each time step in the dataset.

Thanks!



**Jason Brownlee** October 31, 2017 at 2:51 pm #

REPLY ↩

Yes, but you will need to use an extension of ARIMA called ARIMAX. I do not have an example, sorry.



**Dave J** November 5, 2017 at 7:12 am #

REPLY ↩

Hi Dr Brownlee, thanks so much for the tutorials!

I've searched but didn't find anything – perhaps my fault...

But do you have any tutorials or suggestions about forecasting with limited historical observations? Specifically, I'm in a position where some sensors may have a very limited set of historical observations (complete, but short, say it's only been online for a month), but I have many sensors which could possibly be used as historical analogies (multiple years of data).

I've considered constructing a process that uses each large-history sensor as the "Training" set, and iterating over each sensor and finding which sensor best predicts the observed readings for the newer sensors.

However I'm struggling to find any established best practices for this type of thing. Do you have any suggestions for me?

If not I understand, but I really appreciate all the insight you've given over these tutorials and in your book!



**Jason Brownlee** November 6, 2017 at 4:45 am #

REPLY ↩

Great question.

You might be able to use the historical data or models for different but similar sensors (one some dimension). Get creative!



**Dave J** November 6, 2017 at 10:53 am #

REPLY ↩

I would likely just be looking at the RMSE and MAE to gauge accuracy, correct? Is there another measure of fitness I would be wise to consider?



**Jason Brownlee** November 7, 2017 at 9:45 am #

REPLY ↩

No MSE and RMSE are error scores for regression problems. Accuracy is for classification problems (predicting a label).



**Debola** November 11, 2017 at 5:28 am #

REPLY ↩

Hi, Geat tutorial. A question about the difference function. How is it accounting for leap years?

**Jason Brownlee** November 11, 2017 at 9:24 am #

REPLY ↩



It doesn't, that would be a good extension to this tutorial.



**Debola** November 12, 2017 at 12:37 am #

REPLY ↩

Is it possible to apply `seasonal_decompose` on the dataset used in this tutorial since it's a daily forecast. Most applications of `seasonal_decompose` i have seen are usually on monthly and quarterly data



**Jason Brownlee** November 12, 2017 at 9:05 am #

REPLY ↩

Yes, you could use it on this data.



**Akanksha** November 19, 2017 at 4:32 am #

REPLY ↩

Thank you for an amazing tutorial. I wanted to ask if I can store the multiple step values that are predicted in the end of your tutorial into a variable for comparison with actual/real values?



**Jason Brownlee** November 19, 2017 at 11:10 am #

REPLY ↩

Sure, you can assign them to a variable or save them to file.



**Satyajit Pattnaik** December 21, 2017 at 5:01 pm #

REPLY ↩

@Jason, Thanks for this, but my dataset is in a different format, it's in YYYY-MM-DD HH:MI:SS, and the data is hourly data, let say if we have data till 11/25/2017 23:00 5.486691952

And we need to predict the next day's data, so we need to predict our next 24 steps, what needs to be done?

Need a help on this.



**Jason Brownlee** December 22, 2017 at 5:31 am #

REPLY ↩

Sure, you can specify the date-time format when loading the Pandas Series.

You can predict multiple steps using the `predict()` function.



**Satyajit Pattnaik** December 21, 2017 at 8:02 pm #

REPLY ↩

One more question on top of my previous question,  
let say my data is hourly data, and i have one week's data as of now, as per your code do i have to take the days\_in\_year parameter as 7 for my case?

And as per my data's ACF & PACF, my model should be ARIMA(xyz, order=(4,1,2))  
and taking the days\_in\_year parameter as 7, is giving my results, but not sure how correct is that..  
please elaborate a bit @Jason



**Jason Brownlee** December 22, 2017 at 5:32 am #

REPLY ↩

I would recommend tuning the model to your specific data.



**Satyajit Pattnaik** January 3, 2018 at 11:47 pm #

REPLY ↩

Hi Jason,

I am bugging you, but here's my last question, my model is ready and i have predicted the p,d,q values as per the ACF, PACF plots.

Now my code looks like this:

```
1 history = [x for x in train]
2 predictions = list()
3 for t in range(len(test)):
4     model = ARIMA(history, order=(6,1,2))
5     model_fit = model.fit(dispatch=0)
6     output = model_fit.forecast()
7     yhat = output[0]
8     predictions.append(yhat)
9     obs = test[t]
10    history.append(obs)
```

Here, as i am appending obs to the history data, what if i add my prediction to history and then pass it to the model, do i have to run this in a loop to predict pdq values again in a loop?

My question is, if we are doing Recursive multi step forecast do we have to run the history data to multiple ARIMA models, or can we just use history.append(yhat) in the above code and get my results?



**Jason Brownlee** January 4, 2018 at 8:12 am #

REPLY ↩

Recursive multi-step means you will use predictions as history when you re-fit the model.



**Satyajit Pattnaik** January 4, 2018 at 4:48 pm #

REPLY ↩

Reply to my previous response, so predictions to be added as history, that's fine, we will be doing history.append(yhat) instead of history.append(obs), but do we have to run the above code using the same ARIMA model i.e. 6,1,2 or for each history we will determine the pdq values and run on multiple ARIMA models to get the next predictions?



I hope, you are getting my point.

## Leave a Reply

 Name (required) Email (will not be published) (required) Website

SUBMIT COMMENT

## Welcome to Machine Learning Mastery

---



Hi, I'm Dr. Jason Brownlee.

My goal is to make practitioners like YOU awesome at applied machine learning.

[Read More](#)

## Get Good at Time Series Forecasting


---

Need visualizations and forecast models?

Looking for step-by-step tutorials?

Want end-to-end projects?

[Get Started with Time Series Forecasting in Python!](#)



## POPULAR

**Time Series Prediction with LSTM Recurrent Neural Networks in Python with Keras**

JULY 21, 2016

**Your First Machine Learning Project in Python Step-By-Step**

JUNE 10, 2016

**Develop Your First Neural Network in Python With Keras Step-By-Step**

MAY 24, 2016

**Multivariate Time Series Forecasting with LSTMs in Keras**

AUGUST 14, 2017

**How to Setup a Python Environment for Machine Learning and Deep Learning with Anaconda**

MARCH 13, 2017

**Sequence Classification with LSTM Recurrent Neural Networks in Python with Keras**

JULY 26, 2016

**Time Series Forecasting with the Long Short-Term Memory Network in Python**

APRIL 7, 2017

**Regression Tutorial with the Keras Deep Learning Library in Python**

JUNE 9, 2016

**Multi-Class Classification Tutorial with the Keras Deep Learning Library**

JUNE 2, 2016

**How to Grid Search Hyperparameters for Deep Learning Models in Python With Keras**

AUGUST 9, 2016

---

© 2018 Machine Learning Mastery. All Rights Reserved.

[Privacy](#) | [Contact](#) | [About](#)